

# CIS600 - INTRO TO MACHINE LEARNING

## PROJECT REPORT

Gowtham Behara  
Maithreya Pola  
Venkata Krishna Mohan Rallabandi  
Chaitanya Makkapati

## Table of Contents

<b>INTRODUCTION:</b> .....	<b>3</b>
<b>DATA COLLECTION:</b> .....	<b>4</b>
<b>EXPLORATORY DATA ANALYSIS:</b> .....	<b>5</b>
<b>DATA PREPROCESSING</b> .....	<b>7</b>
FEATURE SELECTION .....	7
DATA CLEANING AND MERGING .....	8
<b>MODEL TRAINING:</b> .....	<b>11</b>
<b>MODEL TESTING:</b> .....	<b>13</b>
<b>PREDICTIONS:</b> .....	<b>14</b>
<b>CONCLUSION:</b> .....	<b>18</b>
<b>REFERENCES:</b> .....	<b>19</b>

**INTRODUCTION:**

Regression is a type of supervised learning algorithm in machine learning that is used to predict a continuous value output from a given set of input features. In other words, it is a technique to estimate the relationship between a dependent variable (also known as the target variable) and one or more independent variables (also known as predictor variables or features).

In this project, we're going to deal with a regression problem statement where we're building a machine learning model to predict the delay in flight arrival times. For simplicity, we're just assuming flight from United Airlines with Syracuse (SYR) as its destination. There are 4 origin airports considered, namely, Chicago (ORD), Denver (DEN), Newark (EWR) and Washington (IAD). The flight arrival delays are categorized into 4 different labels: 'EARLY', 'ON-TIME', 'DELAYED' and 'SEVERLY DELAYED'.

We're also considering weather data across all locations since weather plays an important factor in deciding flight arrival times. In fact, weather is one of the leading causes of flight delays and cancellations worldwide. Therefore, airlines and air traffic control centers use advanced weather forecasting technologies and models to predict and monitor weather patterns, assess their potential impact on flight operations, and make necessary adjustments to reduce delays and ensure safety. However, weather remains an unpredictable factor that can disrupt flight schedules and cause inconvenience and frustration for passengers.

The objective of regression analysis is to find the best-fitting line (or curve) that describes the relationship between the input variables, here it's airlines data merged with weather data and the target variable, which is the flight arrival delay. This line or curve can then be used to make predictions on unseen data.

**DATA COLLECTION:**

The first order in any data science/machine learning project is to collect data. Since, we're dealing with predictive analysis, we're going to need historical data to train the model. The model believes that the future depends on historical data, unlike in real world, where there are many other factors influencing the events of the future. In fact, it could be even possible that predictions based on past, and the future are completely different from each other. This brings us to mention a quote from the famous German TV series, *Dark*!

*"Yesterday, today and tomorrow are not consecutive, they are connected in a never-ending circle. Everything is connected."*

For airlines data, we're using [Bureau of Transportation Statistics](#) to fetch historical data of flights arrival times. This will give us information about all United Airlines flights arriving at Syracuse. We're only interested in data from the year 2019 for both airlines and weather data. Once, we download the data as csv, we filter out records that have origin airport as ORD, DEN, EWR and IAD. This will give us required data for factors affecting flight delay due to airlines' causes. Next, we move on to collect weather data.

To collect weather data, we used [Weatherbit](#) APIs to fetch historical weather data. Each member had to use their individual token to fetch this data. We have developed a python script to iteratively fetch hourly weather data from a given location from 2019/1/1 to 2023/3/31. We collected all the historical weather data using the coordinates of all 5 airports, SYR, IAD, ORD, DEN and EWR. The script for this python file is named as weather.py in our project.

We'll also collect forecast weather data using [Weatherbit Forecast API](#) for the duration between 2023/04/21 to 2023/04/24 and we'll use this data as features in prediction dataset. This will be covered in the PREDICTIONS section of the report.


After this step, we now have historical airlines data and historical weather data for all 5 airports since the beginning of 2019. The next step involves understanding and refining of the data for better model training.

## EXPLORATORY DATA ANALYSIS:


Exploratory Data Analysis (EDA) is a process of analyzing and understanding a dataset to gain insights and make informed decisions. It involves a variety of techniques to investigate the dataset and uncover patterns, anomalies, and relationships between variables.

Some common techniques used in EDA include data visualization, summary statistics, correlation analysis, and hypothesis testing. Data visualization is particularly useful for EDA because it can quickly reveal patterns and relationships in the data that may be difficult to discern from tables or summary statistics alone.


The following heatmap is generated to understand correlation between columns present in train data. We can see that 'Arrival Delay' is more correlated with other factors like 'Delay Late Aircraft Arrival', 'Delay Weather' and 'Delay Carrier'. Hence, we're going to retain these features. The 'Delay Security (Minutes)' features is nowhere related to any other feature because it has all values as zero. We are going to neglect all such columns.



In this step, we're visualizing the temperature distribution of Syracuse through histogram. We can see that most of the temperatures in Syracuse are ranging from 0-5C which is expected. Also, the distribution shows that the data is very consistent throughout the dataset.



Here, we're looking at Dew Point frequency distribution in New Jersey (EWR) through histogram. We can confirm that there's much dew present in this location by verifying that higher dewpoint values have more frequency.



Similarly, we can check for other frequency distributions of various features to assess the data quality and understand their patterns.

## DATA PREPROCESSING

### FEATURE SELECTION

Data preprocessing is a critical step in machine learning that involves preparing raw data to be used in a machine learning model. The main goal of data preprocessing is to transform the data into a format that can be easily understood and used by machine learning algorithms.

Data preprocessing typically involves several steps, including data cleaning, data transformation, and data reduction. Data cleaning involves removing or correcting any missing, incorrect, or irrelevant data. Data transformation involves converting the data into a format that can be more easily analyzed, such as scaling or encoding categorical variables. Data reduction involves selecting a subset of the available data, such as by removing redundant features or instances to reduce the complexity of the model.

During the data pre-processing stage in this project, we are carefully selecting the relevant features from both the weather data and airline data. This involves identifying the key variables or attributes that are likely to have a significant impact on the outcome of our analysis, while discarding any extraneous or irrelevant data.

In the case of weather data, we might consider variables such as local timestamp, temp, clouds, rh, dew point, precipitation, snow, visibility. These factors can have a significant impact on flight schedules and delays, and are therefore important to consider when analyzing airline data.

```
weather_factors = ['timestamp_local', 'temp', 'clouds', 'rh', 'dewpt', 'precip', 'snow', 'vis']  
  
syr_weather_data = syr_weather_data[weather_factors]  
ord_weather_data = ord_weather_data[weather_factors]  
den_weather_data = den_weather_data[weather_factors]  
ewr_weather_data = ewr_weather_data[weather_factors]  
iad_weather_data = iad_weather_data[weather_factors]
```

When it comes to airline data, we may be interested in variables such as Date, Delay National Aviation System (Minutes), Delay Weather (Minutes), Delay Carrier (Minutes), Delay Late Aircraft Arrival (Minutes), Arrival Delay (Minutes), Scheduled Arrival Time, Origin Airport. By selecting the most relevant features from both the weather and airline datasets, we can ensure that our analysis is based on accurate and meaningful data, ultimately leading to more informed decision-making.

```

airlines_factors = ['Date (MM/DD/YYYY)',
                    'Delay National Aviation System (Minutes)',
                    'Delay Weather (Minutes)',
                    'Delay Carrier (Minutes)',
                    'Delay Late Aircraft Arrival (Minutes)',
                    'Arrival Delay (Minutes)',
                    'Scheduled Arrival Time',
                    'Origin Airport']

airlines_data = airlines_data[airlines_factors]
airlines_data = airlines_data.rename(columns={'Date (MM/DD/YYYY)': 'Date',
                                             'Delay Carrier (Minutes)': 'Delay Carrier',
                                             'Delay Weather (Minutes)': 'Delay Weather',
                                             'Delay Late Aircraft Arrival (Minutes)': 'Delay Late Aircraft Arrival',
                                             'Arrival Delay (Minutes)': 'Arrival Delay',
                                             'Delay National Aviation System (Minutes)': 'Delay National Aviation Syst

airlines_data.columns

syr_weather_data = syr_weather_data.rename(columns={'timestamp_local': 'time'})
syr_weather_data.columns

ord_weather_data = ord_weather_data.rename(columns={'timestamp_local': 'time'})
ord_weather_data.columns

den_weather_data = den_weather_data.rename(columns={'timestamp_local': 'time'})
den_weather_data.columns

ewr_weather_data = ewr_weather_data.rename(columns={'timestamp_local': 'time'})
ewr_weather_data.columns

iad_weather_data = iad_weather_data.rename(columns={'timestamp_local': 'time'})
iad_weather_data.columns

```

## DATA CLEANING AND MERGING

Data merging is the process of combining multiple datasets into a single dataset. In machine learning, data merging can be used to enrich or supplement existing data, as well as to prepare data for training and testing machine learning models.

There are several methods for merging datasets, including concatenation, merging, and joining. Concatenation is used to combine datasets that have the same structure and variables, but different instances or samples. Merging is used to combine datasets that have different variables, but share common instances or samples. Joining is used to combine datasets that have common variables, but different instances or samples.

The 'airlines\_data' dataset undergoes processing whereby a new column called 'time' is generated through the combination of the 'Date' and 'Scheduled Arrival Time' columns. This new column is converted into a datetime object and rounded to the nearest hour. After that, the 'Date' and 'Scheduled Arrival Time' columns are removed from the dataset. To confirm that the 'time' column has been successfully added, the modified dataset is displayed.



```
airlines_data['time'] = pd.to_datetime(airlines_data['Date'].astype(str) + 'T' +
                                     airlines_data['Scheduled Arrival Time'].astype(str)).round('H')
airlines_data = airlines_data.drop(columns=['Date', 'Scheduled Arrival Time'])
airlines_data.head()
```

We combined 'airlines\_data' and 'syr\_weather\_data' datasets by using the 'time' column as the reference to match the data. The output dataset includes all columns from both datasets, and only the matching rows are included. We used a 'left' join type where all rows from 'airlines\_data' are included, and any missing values in 'syr\_weather\_data' are replaced with NaN values. This merged dataset can be analyzed to explore possible correlations between weather conditions and flight delays.

```
airlines_data_cp = airlines_data.copy()
syr_dataset = airlines_data_cp.merge(syr_weather_data, on='time', how='left')
```

Our analysis involves examining the connection between weather conditions and flight delays across all five airports. We rename the 'timestamp\_local' column in 'ord\_weather\_data' to 'time'.

- Initially, we use the pandas library in Python to manipulate data and carry out correlation analysis on the 'airlines\_data' and 'ord\_weather\_data' datasets.
- We then merge 'ord\_weather\_data' and 'ord\_dataset' datasets on the 'time' column using a 'left' join type to create a new dataset called 'ord\_dataset'.
- We calculate the correlation matrix for the 'ord\_dataset' dataset using the corr() method.

```
corrmat = syr_dataset.corr()
corrmat['Arrival Delay'].sort_values()
```

- Next, we select the 'Arrival Delay' column and use the 'sort\_values()' method to sort the correlation values in ascending order, identifying variables that have the strongest correlation with arrival delay at each airport.

```
ord_dataset = airlines_data[airlines_data['Origin Airport'] == 'ORD']
ord_dataset.columns

ord_weather_data = ord_weather_data.rename(columns={'timestamp_local': 'time'})
ord_weather_data.columns

ord_dataset = ord_dataset.merge(ord_weather_data, on='time', how='left')

ord_dataset.columns

ord_dataset.head()

corrmat = ord_dataset.corr()
corrmat['Arrival Delay'].sort_values()
```

```
den_dataset = airlines_data[airlines_data['Origin Airport'] == 'DEN']
den_dataset.shape

den_weather_data = den_weather_data.rename(columns={'timestamp_local': 'time'})
den_weather_data.columns

den_dataset = den_dataset.merge(den_weather_data, on='time', how='left')

den_dataset.columns

den_dataset.head()

corrmat = den_dataset.corr()
corrmat['Arrival Delay'].sort_values()
```

```
ewr_dataset = airlines_data[airlines_data['Origin Airport'] == 'EWR']
ewr_dataset.shape

ewr_weather_data = ewr_weather_data.rename(columns={'timestamp_local': 'time'})
ewr_weather_data.columns

ewr_dataset = ewr_dataset.merge(ewr_weather_data, on='time', how='left')

ewr_dataset.columns

ewr_dataset.head()

corrmat = ewr_dataset.corr()
corrmat['Arrival Delay'].sort_values()
```

```
iad_dataset = airlines_data[airlines_data['Origin Airport'] == 'IAD']
iad_dataset.shape

iad_weather_data = iad_weather_data.rename(columns={'timestamp_local': 'time'})
iad_weather_data.columns

iad_dataset = iad_dataset.merge(iad_weather_data, on='time', how='left')

iad_dataset.columns

iad_dataset.head()

corrmat = iad_dataset.corr()
corrmat['Arrival Delay'].sort_values()
```

- This analysis provides insight into factors contributing to flight delays, enabling airlines and airports to take necessary actions to minimize them.
- Lastly, we merge all datasets for each airport and shuffle the dataset.

```
dataset = pd.concat([syr_dataset, ord_dataset, iad_dataset, den_dataset, ewr_dataset])  
  
from sklearn.utils import shuffle  
dataset = shuffle(dataset)
```

We remove the 'time' column from the 'dataset' dataframe and create dummy variables for the 'Origin Airport' column using the `get_dummies()` method, which converts categorical variables into binary indicator variables.

```
dataset = dataset.drop(columns=['time'])  
dataset = pd.get_dummies(dataset, columns=['Origin Airport'], drop_first=True)  
dataset.shape
```

## MODEL TRAINING:

We've finally arrived at the most interesting part of our project, i.e., model building/training. Now, we are standing with 3000+ rows of data with 14 columns as features. Note that, all the feature values are numerical in nature after all the preprocessing has been done in the previous step. This is particularly advantageous in building a regression model.

The 'Arrival Delay' column is selected as the Y variable for our model. This is the numerical delay in arrival time for our flights. We're going to estimate values for this column in our prediction dataset. All the remaining columns are part of the (X) feature matrix.

We're going to split our dataset into train and test set (20%) with random state of 42. Setting the random state to 42 is a common convention in the machine learning community, as it is a reference to "The Hitchhiker's Guide to the Galaxy" by Douglas Adams, where the number 42 is famously described as the "Answer to the Ultimate Question of Life, the Universe, and Everything." While this reference is mostly meant to be humorous, it has become a popular choice for setting the random state in machine learning algorithms.

```
In [24]: X_train, X_test, y_train, y_test = train_test_split(dataset.drop(columns=['Arrival Delay']), dataset['Arrival Delay'],
                                                         test_size=0.2, random_state=42)

X_train
X_test
y_train
y_test
```

Next, we proceed to model training by using the ‘DecisionTreeRegressor’ from `sklearn.tree` module. `DecisionTreeRegressor` is a type of regression algorithm used in machine learning. It is based on the decision tree model, which involves recursively partitioning the input space into smaller regions based on the values of the input features, and fitting a simple model (e.g., a constant value) to each region.

The `DecisionTreeRegressor` algorithm works by repeatedly splitting the dataset into subsets based on the values of a single input feature. The splitting criterion is chosen to maximize the reduction in the variance of the target variable (or some other measure of homogeneity), and the process continues until a stopping criterion is met (e.g., a maximum depth of the tree or a minimum number of samples per leaf).

One advantage of `DecisionTreeRegressor` is that it is relatively simple and easy to interpret, as the resulting tree structure can be visualized and analyzed to gain insights into the relationship between the input features and the target variable. However, it is prone to overfitting, especially when the tree is allowed to grow too deep and may not generalize well to unseen data. Therefore, various techniques such as pruning, ensemble methods, and regularization can be used to improve the performance of `DecisionTreeRegressor` and reduce overfitting.

The following is the code snippet for implementing model training:

```
In [39]: from sklearn.tree import DecisionTreeRegressor

clf = DecisionTreeRegressor(random_state=50)

clf = clf.fit(X_train_scaled, y_train)

clf.score(X_train_scaled, y_train)
```

```
Out[39]: 0.9991997279777147
```

We’re also computing the  $R^2$  score of the trained model on the train data itself to understand how much the model has learnt from the train data. Looks like the model has very much understood its assignment. >99% score implies a very good train data score. But we need to worry about the test data score in order to understand if overfitting has occurred or not.

**MODEL TESTING:**

In this step, we're going to test the accuracy of our model with the test dataset. Although this step looks simple, it's very necessary to evaluate the accuracy of our model on unseen data. If the test score is comparatively very less than our train data score, then we need to retrain our model with better data preprocessing, hyper-parameters, more training data, etc.

We're going to use predict function from our trained model and generate a test\_output dataframe with column name as 'pred\_Arrival\_Delay'.

Next, we're going to merge with actual y\_test and compute the mean\_absolute\_error to get an estimate. An m\_a\_e of 15.36 looks decent. Also, the ratio of mean\_absolute\_error to mean of actual arrival delay is 1.34 which looks good enough.

```
In [41]: X_test_scaled = pd.DataFrame(sc.transform(X_test), columns = X_test.columns, index = X_test.index)
```

```
In [42]: test_output = pd.DataFrame(clf.predict(X_test_scaled), index = X_test_scaled.index, columns = ['pred_Arrival_Delay'])
test_output.head()
```

```
Out[42]:
```

	pred_Arrival_Delay
1841	-15.00
253	50.00
286	-30.00
1013	14.00
1729	88.00

```
In [43]: test_output = test_output.merge(y_test, left_index = True, right_index = True)
test_output.head()
mae = abs(test_output['pred_Arrival_Delay'] - test_output['Arrival Delay']).mean()
mae.round(2)
```

```
Out[43]:
```

	pred_Arrival_Delay	Arrival Delay
0	51.00	51
3	13.00	-12
3	13.00	-4
3	3.00	-12
3	3.00	-4

```
Out[43]: 15.36
```

```
In [44]: (abs(test_output['pred_Arrival_Delay'] - test_output['Arrival Delay']).mean()/test_output['Arrival Delay'].mean()).round(2)
```

```
Out[44]: 1.34
```

```
In [45]: clf.score(X_test_scaled, y_test)
```

```
Out[45]: 0.8190666349281132
```

Finally, the model score on test data is around 82%. Now, this is close to our train data score which is >99%. Hence, it is convincible to say that our model has not been overfit on the train data.

**PREDICTIONS:**

```
import json

with open('../datasets/project/hs/forecast/syr-forecast.json', 'r') as f:
    syr_forecast_data = json.load(f)
    syr_forecast_df = pd.json_normalize(syr_forecast_data['data'])
    syr_forecast_df = syr_forecast_df[weather_factors]
    syr_forecast_df = syr_forecast_df.rename(columns={'timestamp_local': 'time'})
    syr_forecast_df['time'] = syr_forecast_df['time'].astype(np.datetime64)
    syr_forecast_df.columns
```

It reads a JSON file containing forecast data for a specific location and loads it into a Pandas dataframe. The dataframe is then preprocessed to only keep the relevant features related to weather factors. The 'timestamp\_local' feature is renamed to 'time', and the 'time' feature is converted to a datetime format.

```
with open('../datasets/project/hs/forecast/den-forecast.json', 'r') as f:
    den_forecast_data = json.load(f)
    den_forecast_df = pd.json_normalize(den_forecast_data['data'])
    den_forecast_df = den_forecast_df[weather_factors]
    den_forecast_df = den_forecast_df.rename(columns={'timestamp_local': 'time'})
    den_forecast_df['time'] = syr_forecast_df['time'].astype(np.datetime64)
    den_forecast_df.columns

with open('../datasets/project/hs/forecast/ewr-forecast.json', 'r') as f:
    ewr_forecast_data = json.load(f)
    ewr_forecast_df = pd.json_normalize(ewr_forecast_data['data'])
    ewr_forecast_df = ewr_forecast_df[weather_factors]
    ewr_forecast_df = ewr_forecast_df.rename(columns={'timestamp_local': 'time'})
    ewr_forecast_df['time'] = syr_forecast_df['time'].astype(np.datetime64)
    ewr_forecast_df.columns

with open('../datasets/project/hs/forecast/iad-forecast.json', 'r') as f:
    iad_forecast_data = json.load(f)
    iad_forecast_df = pd.json_normalize(iad_forecast_data['data'])
    iad_forecast_df = iad_forecast_df[weather_factors]
    iad_forecast_df = iad_forecast_df.rename(columns={'timestamp_local': 'time'})
    iad_forecast_df['time'] = syr_forecast_df['time'].astype(np.datetime64)
    iad_forecast_df.columns

with open('../datasets/project/hs/forecast/ord-forecast.json', 'r') as f:
    ord_forecast_data = json.load(f)
    ord_forecast_df = pd.json_normalize(ord_forecast_data['data'])
    ord_forecast_df = ord_forecast_df[weather_factors]
    ord_forecast_df = ord_forecast_df.rename(columns={'timestamp_local': 'time'})
    ord_forecast_df['time'] = syr_forecast_df['time'].astype(np.datetime64)
    ord_forecast_df.columns
```

The same process is being used to load and preprocess data for different locations. The code reads forecast data for different airports from JSON files and converts it into a Pandas dataframe using the `pd.json_normalize()` function. Finally, the column names are checked to ensure that they are in the correct format. By repeating this process for each location, we can prepare the data for the machine learning model to make predictions for different airports. This preprocessed data can then be used as input to a trained machine learning model to make predictions about future weather conditions for that location.

```
pred_data['time'] = pd.to_datetime(pred_data['Date'].astype(str) + 'T' +  
                                   pred_data['Arrival Time'].astype(str)).round('H')  
pred_data = pred_data.drop(columns=['Date', 'Arrival Time', 'Day', 'Flight Number'])  
pred_data.head()
```

The above code creates a new 'time' column by combining the 'Date' and 'Arrival Time' columns, which are then converted to datetime objects and rounded to the nearest hour. These new datetime features can be used to capture potential patterns or relationships between arrival times and other features, which can help to improve the accuracy of the machine learning model. Additionally, unnecessary columns are dropped using the 'drop' function, which reduces the dimensionality of the data and may improve model efficiency. Overall, these preprocessing steps can help to improve the quality of the data and increase the performance of the machine learning model.

```
ord_pred = pred_data[pred_data['Origin Airport'] == 'ORD']  
ord_pred.shape  
  
ord_pred = ord_pred.merge(ord_forecast_df, on='time', how='left')  
ord_pred.head()
```

The code above first filters the predicted data based on the Origin Airport value of "ORD". This filtered data is then merged with the ORD airport weather forecast data, which was previously loaded and preprocessed. The merging is performed based on a common column named "time". This will allow the predicted data to be augmented with additional information from the forecast dataset. The resulting merged data can be further processed and used as input to a machine learning model for prediction or analysis.



```
den_pred = pred_data[pred_data['Origin Airport'] == 'DEN']
den_pred.shape

den_pred = den_pred.merge(den_forecast_df, on='time', how='left')
den_pred.head()

ewr_pred = pred_data[pred_data['Origin Airport'] == 'EWR']
ewr_pred.shape

ewr_pred = ewr_pred.merge(ewr_forecast_df, on='time', how='left')
ewr_pred.head()

iad_pred = pred_data[pred_data['Origin Airport'] == 'IAD']
iad_pred.shape

iad_pred = iad_pred.merge(iad_forecast_df, on='time', how='left')
iad_pred.head()
```

The above code is a repetitive process that performs the same operation on multiple airports. For each airport, the code filters the prediction data to only include flights that depart from that airport. It then merges the filtered prediction data with the respective airport's weather forecast data on the 'time' column using a left join. The resulting dataframe contains all the flight data along with the corresponding weather forecast data for that airport. This process is repeated for multiple airports including DEN, EWR, IAD, and ORD.

```
d_nas = dataset['Delay National Aviation System'].mean()
d_w = dataset['Delay Weather'].mean()
d_c = dataset['Delay Carrier'].mean()
d_laa = dataset['Delay Late Aircraft Arrival'].mean()
```

The above code calculates the mean of four delay types in a given dataset: National Aviation System (NAS) delay, Weather delay, Carrier delay, and Late Aircraft Arrival (LAA) delay. These delay types can be used as features in a regression model to predict flight delays. By calculating their means, we can gain insight into the average delay times for each type of delay, which can help us understand which type of delay has the most impact on flight delays overall.

```
pred_dataset = pd.concat([ord_pred, den_pred, ewr_pred, iad_pred])
```

In this step of the machine learning process, separate data frames are created for each airport location: ORD, DEN, EWR, and IAD. These data frames contain the relevant input features such as temperature, precipitation, and cloud cover, as well as the mean delay times for each delay category. After the creation of these data frames, they are concatenated into a single data frame, called 'pred\_dataset'. This data frame will be used as input for the predictive model in the next steps of the machine learning pipeline.



```
pred_dataset.insert(2, 'Delay National Aviation System', d_nas)
pred_dataset.insert(2, 'Delay Weather', d_w)
pred_dataset.insert(2, 'Delay Carrier', d_c)
pred_dataset.insert(2, 'Delay Late Aircraft Arrival', d_laa)
pred_dataset.head()
```

Here, the **pred\_dataset** is created by concatenating four datasets for different airports. The mean delays for different factors such as National Aviation System, weather, carrier, and late aircraft arrival are calculated from the original dataset **dataset**. These mean values are then inserted as new columns into the **pred\_dataset** using the **insert()** method. By adding these mean delay columns to the prediction dataset, the model can use them as additional features to make more accurate predictions.

```
ordered_cols = ['Delay National Aviation System', 'Delay Weather', 'Delay Carrier', 'Delay Late Aircraft Arrival',
                'temp', 'clouds', 'rh', 'dewpt', 'precip', 'snow', 'vis',
                'Origin Airport_EWR', 'Origin Airport_IAD', 'Origin Airport_ORD']
pred_dataset = pred_dataset.reindex(columns=ordered_cols)
```

We are reordering the columns of the **pred\_dataset** dataframe for the ease of model training. The first four columns are related to the mean delays, followed by columns that contain weather factors such as temperature, cloudiness, relative humidity, dew point, precipitation, snow, and visibility. The next three columns are binary encoded airport codes representing the origin airport of the flight. The **reindex()** function is used to rearrange the columns in the desired order specified by the **ordered\_cols** list. This step helps in maintaining the consistency of the input data during the model training process.

```
pred_output = pd.DataFrame(clf.predict(pred_dataset), index = pred_dataset.index, columns = ['Pred Arrival Delay'])
pred_output.head()
```

Here a prediction is made using the trained machine learning model, **clf**, to predict the arrival delay of flights. The **pred\_dataset**, which contains the relevant features required for the prediction, is passed into the **clf** model, and the predicted arrival delay is returned as a Pandas DataFrame, **pred\_output**. The index of the DataFrame is set to **pred\_dataset.index**, which is the index of the input DataFrame, and the column is named 'Pred Arrival Delay'. The resulting **pred\_output** DataFrame contains the predicted arrival delay for each flight in the **pred\_dataset**.

```
bins = [-float('inf'), -10, 10, 30, float('inf')]
labels = ['Early', 'On-time', 'Late', 'Severly Late']

pred_output['Status (Early, On-time, Late, Severly Late)'] = pd.cut(pred_output['Pred Arrival Delay'],
                                                                    bins=bins, labels=labels)
pred_output.head()
```

Binning is a process of transforming continuous numerical features into categorical features. In this code block, we are creating bins for the predicted arrival delays. We specify four bins with the labels 'Early', 'On-time', 'Late', 'Severly Late' and use the `pd.cut()` function to assign each predicted arrival delay to a bin. The bin intervals are defined as -inf to -10, -10 to 10, 10 to 30, and 30 to inf. The resulting output is a new column in the `pred_output` dataframe containing the categorical labels of the predicted arrival delays. This process can help in better understanding and visualizing the predicted arrival delays.

```
pred_output = pred_output.drop(columns=['Pred Arrival Delay'])
```

```
pred_output.to_csv('project csv(Apr 21-24).csv', index=False)
```

The code above drops the "Pred Arrival Delay" column from the "pred\_output" dataframe, which contains the numerical predicted arrival delay for each flight. We only retain the converted labelled column 'Status (Early, On-time, Late, Severly Late)' in our final output dataframe. Finally, we're saving this dataframe into a csv as the output file of our project.

## CONCLUSION:

In conclusion, the regression analysis using decision trees has proven to be a useful tool for predicting the arrival time of flights into Syracuse (SYR) from four major airports - Chicago (ORD), Denver (DEN), Newark (EWR), and Washington (IAD) - up to four days in advance. The results of the analysis show that the model was able to accurately predict whether flights would arrive early, on-time, delayed, or severely delayed, based on a variety of factors such as weather conditions, flight schedules, and airport congestion.

The decision tree algorithm was particularly effective in identifying the most important variables for predicting flight arrival times, allowing for a more accurate and efficient prediction model. The model was able to provide valuable insights to airlines, airports, and travelers, allowing them to plan and adjust their schedules accordingly.

Overall, the use of regression analysis using decision trees has the potential to significantly improve the efficiency and reliability of the airline industry, ultimately benefiting both businesses and consumers. With further research and development, this approach could potentially be applied to other areas of transportation and logistics as well.

**REFERENCES:**

- <https://www.transtats.bts.gov/ontime/arrivals.aspx>
- <https://www.weatherbit.io/>
- <https://realpython.com/pandas-merge-join-and-concat/>
- A review paper on data preprocessing: A critical phase in web usage mining process.  
[Sanjay Kumar Dwivedi](#), [Bhupesh Rawat](#)